

# Advanced SQL

---

03 — Standard and Non-Standard Data Types

Summer 2020

Torsten Grust  
Universität Tübingen, Germany

# 1 | Data Types in (Postgre)SQL

---

- The set of supported **data types** in PostgreSQL is varied:<sup>1</sup>

```
SELECT string_agg(t.typname, ' ') AS "data types"
FROM   pg_catalog.pg_type AS t
WHERE  t.typelem = 0 AND t.typrelid = 0;
```

data types										
bool	bytea	char	int8	int2	int4	regproc	text	oid	tid	C R
oid	tid	xid	cid	json	xml	pg_node_tree	pg_ddl_command			C R
smgr	path	polygon	float4	float8	abstime	reltime				C R
tinterval	unknown	circle	money	macaddr	inet	cidr	...			


<sup>1</sup> See <https://www.postgresql.org/docs/current/datatype.html>

## 2 | SQL Type Casts

---

Convert type of value  $e$  to  $\tau$  at *runtime* via a **type cast**:

<code>CAST (&lt;e&gt; AS &lt;τ&gt;)</code> <code>&lt;e&gt; :: &lt;τ&gt;</code> <code>&lt;τ&gt;(&lt;e&gt;)</code>	$\left. \vphantom{\begin{array}{l} \text{CAST} \\ \text{::} \\ \text{func} \end{array}} \right\} \text{equivalent}$	(SQL standard) (PostgreSQLism, cf. FP) (if $\tau$ valid func name)
--	---	--

-  Type casts can fail at runtime.
- SQL performs **implicit casts** when the required target type is unambiguous (e.g. on insertion into a table column):

```
INSERT INTO T(a,b,c,d) VALUES (6.2, NULL, 'true', '0')
```






-- implicitly casts to: int text boolean int

## Literals (Casts From **text**)

---

SQL supports **literal syntax** for dozens of data types in terms of **casts from type **text****:

<code>CAST ('&lt;literal&gt;' AS &lt;τ&gt;)</code> <code>'&lt;literal&gt;' :: &lt;τ&gt;</code> <code>&lt;τ&gt; '&lt;literal&gt;'</code>	$\left. \vphantom{\begin{array}{l} \text{CAST ('<literal>' AS <τ>)} \\ \text{'<literal>' :: <τ>} \\ \text{<τ> '<literal>'}$	succeeds if <i>&lt;literal&gt;</i> has a valid interpretation as <i>&lt;τ&gt;</i> (without cast $\Rightarrow$ type <u>text</u> )
---	---	--

- Embed complex literals (e.g., dates/times, JSON, XML, geometric objects) in SQL source.
- Casts from **text** to **τ** attempted **implicitly** if target type **τ** known. Also vital when importing data from text or CSV files (*input conversion*).

### 3 | Text Data Types

---

<u>char</u>	-- $\equiv$ char(1)
<u>char</u> (< <i>n</i> >)	-- fixed length, blank ( ) padded if needed
<u>varchar</u> (< <i>n</i> >)	-- varying length $\leq n$ characters
<u>text</u>	-- varying length, unlimited

- Length limits measured in characters, *not* bytes.  
(PostgreSQL: max size  $\approx$  1 Gb. Large text is “TOASTed.”)
- For `char(<n>)`, `varchar(<n>)` length limits are enforced:
  1. Excess characters (other than ) yield runtime errors.
  2. Explicit casts truncate to length *n*.
- `char(<n>)` always *printed/stored* using *n* characters: pad with ). ⚠ Trailing blanks removed before computation.

## 4 : **NUMERIC**:<sup>2</sup> Large Numeric Values with Exact Arithmetics

---

	<i>scale</i>
	┐
<u>numeric</u> ( <i>&lt;precision&gt;</i> , <i>&lt;scale&gt;</i> )	1234567.890
	└────────────────┘
	<i>precision</i> (# of digits)

- Shorthand: `numeric(<precision>,0)`  $\equiv$  `numeric(<precision>)`.  
`numeric`  $\equiv$  " $\infty$  precision" (PostgreSQL limit: 100000+).
- Exact arithmetics, but computationally heavy.
- Leading/trailing 0s *not* stored  $\Rightarrow$  variable-length data.

<sup>2</sup> Synonymous: `decimal`.

## Long **NUMERIC**s Carry a Lot(!) of Bits (Tupper's Formula)

---

A **numeric** value of hundreds of digits can encode a lot of information in a single table cell. Consider:

**Tupper's formula** (with  $x \in [0, 106)$  and  $y \in [k, k+17)$ )


$$\frac{1}{2} < \lfloor \text{mod}(\lfloor y/17 \rfloor 2^{-17} \lfloor x \rfloor - \text{mod}(\lfloor y \rfloor, 17), 2) \rfloor$$

decodes  $k = 9609397 \dots \langle 530 \text{ digits omitted} \rangle \dots 8404719$  to give:

$$\frac{1}{2} < \lfloor \text{mod}(\lfloor y/17 \rfloor 2^{-17} \lfloor x \rfloor - \text{mod}(\lfloor y \rfloor, 17), 2) \rfloor$$

## 5 | Timestamps and Time Intervals

---




- Types: `timestamp`  $\equiv$  (`date`, `time`). Casts between types: `timestamp`  $\rightarrow$  `time/date`  $\checkmark$ , `date`  $\rightarrow$  `timestamp` assumes 00:00:00. Optional timezone support: `< $\tau$ >` with `time zone` or `< $\tau$ >tz`.
- Type `interval` represents timestamp differences.
- Resolution: `timestamp/time/interval`: 1  $\mu$ s, `date`: 1 day.



## Date/Time Literals: PostgreSQL

---

- Literal input and output: flexible/human-readable ,  
affected by `SET datestyle='{German,ISO},{MDY,DMY,YMD}'`  

$\underbrace{\hspace{10em}}$   
output

$\underbrace{\hspace{10em}}$   
input
- `timestamp` literal  $\equiv$  `'<date literal>_<time literal>'`
- `interval` literal (fields optional, `s` may be fractional)  $\equiv$   
`'<n>years <n>months <n>days <n>hours <n>mins <s>secs'`
- Special literals:
  - `timestamp`: `'epoch'`, `'[-]infinity'`, `'now'`
  - `date`: `'today'`, `'yesterday'`, `'tomorrow'`, `'now'`

## Computing with Time

---

- Timestamp arithmetic via `+`, `-` (`interval` also `*`, `/`):

```
SELECT ('now'::timestamp - 'yesterday'::date)::interval
```

interval
1 day 17:27:47.454803

- PostgreSQL: *Extensive* library of date/time functions including:
  - `timeofday()` (⚠ yields `text`)
  - `extract(<field> from .)`
  - `make_date(.,.,.)`, `make_time(...)`, `make_timestamp(...)`
  - comparisons (`=`, `<`, `...`), `(.,.) overlaps (.,.)`

## 6 | Enumerations

---

Create a *new* type  $\tau$ , incomparable with any other. Explicitly **enumerate** the literals  $v_i$  of  $\tau$ :

```
CREATE TYPE < $\tau$ > AS ENUM (< $v_1$ >, ..., < $v_n$ >);
```

```
SELECT < $v_i$ >::< $\tau$ >;
```

- Literals  $v_i$  in case-sensitive string notation '...'.  
(Storage: 4 bytes, regardless of literal length.)
- Implicit ordering:  $v_i < v_j$  (aggregates **MIN**, **MAX** ✓).

## 7 : Bit Strings

---

- Data type `bit(<n>)` stores strings of `n` binary digits (storage: 1 byte per 8 bits + constant small overhead).
- Literals:



```
SELECT B'00101010', X'2A', '00101010'::bit(8), 42::bit(8)
```

$\underbrace{\hspace{10em}}$   
 $2 \times 4 \text{ bits}$

- Bitwise operations: `&` (and), `|` (or), `#` (xor), `~` (not), `<</>>` (shift left/right), `get_bit(.,.)`, `set_bit(.,.)`
- String-like operations: `||` (concatenation), `length(.)`, `bit_length(.)`, `position(. in .)`, ...

## 8 : Binary Arrays (BLOBs)

---

Store **binary large object blocks** (BLOBs; ,  in column *B* of type *bytea*) in-line with alpha-numeric data. BLOBs remain *uninterpreted* by DBMS:



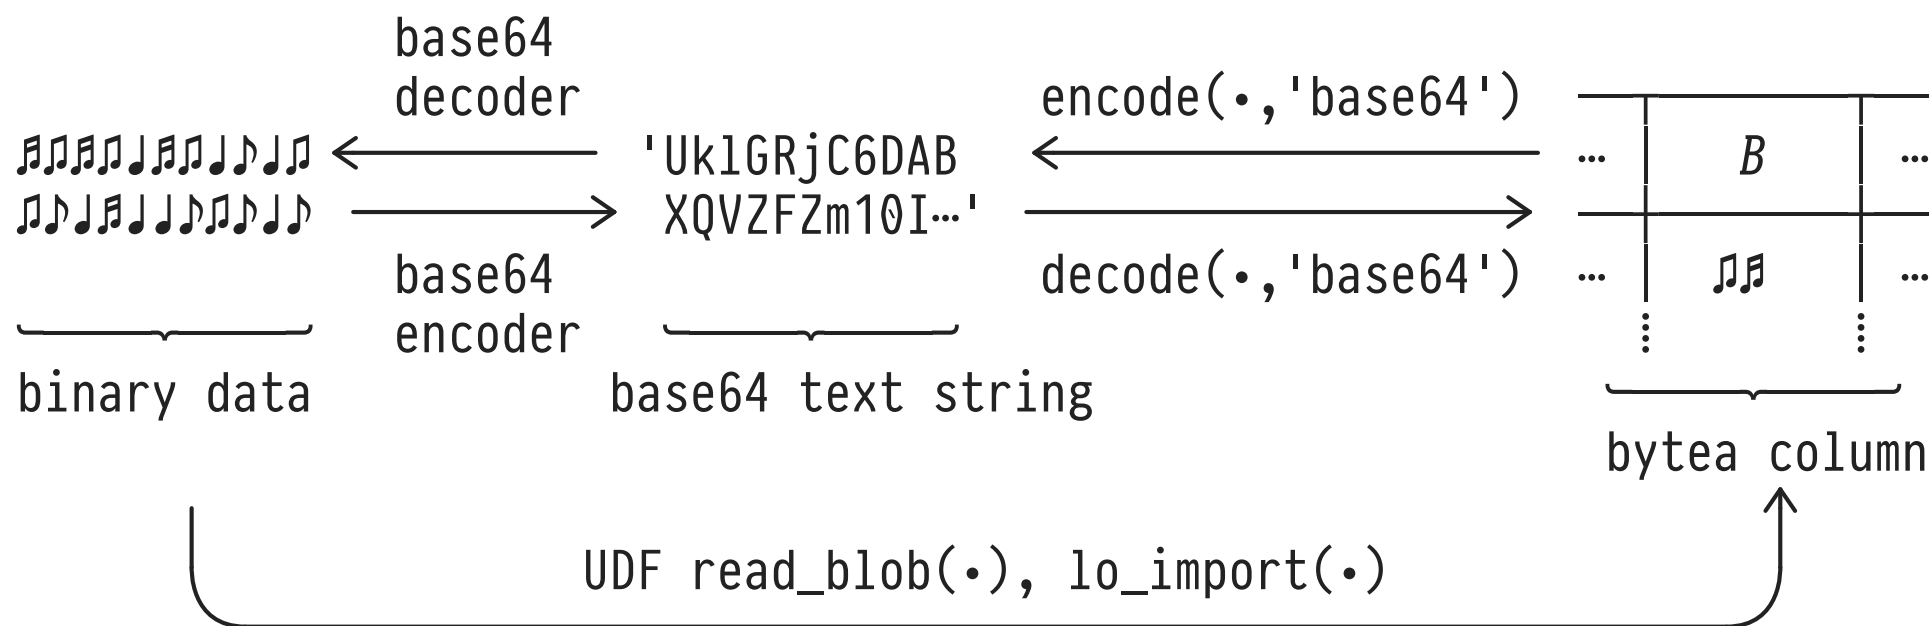
...	<i>K</i>	<i>B</i> :: <i>bytea</i>	<i>P</i>	...
	⋮	⋮	⋮	
	$k_i$		$p_i$	
	$k_j$		$p_j$	
	⋮	⋮	⋮	

Table *T*

- Typical setup:
  - BLOBs stored alongside identifying **key** data (column *K*).
  - Additional **properties** (meta data, column(s) *P*) made explicit to filter/group/order BLOBs.

## Encoding/Decoding BLOBs

- Import/export `bytea` data via textual encoding (e.g., base64) or directly from/to binary files:



⚠ File I/O performed by DBMS server (paths, permissions).

## 9 | Ranges (Intervals)

---


Given lower and/or upper bounds  $\ell$ ,  $u$  of an ordered type  $\langle\tau\rangle \in \{\text{int4}, \text{int8}, \text{num}(\text{eric}), \text{timestamp}, \text{date}\}$ , construct **range** literals of type  $\langle\tau\rangle\text{range}$  via

$[\langle\ell\rangle, \langle u\rangle]$	$\ell \leq x \leq u$	$[\text{-----}]$
$[\langle\ell\rangle, \langle u\rangle)$	$\ell \leq x < u$	$[\text{-----}[$
$(\text{ }, \langle u\rangle]$	$x \leq u$	$[\text{-----}]$
$(\langle\ell\rangle, \text{ })$	$\ell < x$	$[\text{-----}]$
empty	$\emptyset$	$[\text{-----}]$

- Alternatively use function  $\langle\tau\rangle\text{range}(\langle\ell\rangle, \langle u\rangle, '[]')$ , **NULL** represents no bound ( $\infty$ ).

# Range Operations

---




$r_1 @> p$	$r_3 <@ r_1$	contains, contained by
$r_1 - - r_2$		is adjacent to
$r_3 << r_2$	$r_1 << r_2$	strictly left of
$r_2 + r_3$		union
$r_1 * r_3$		intersection
$r_1 \&\& r_3$		overlaps

- Additional family of range-supporting functions:
  - `lower(·)`, `upper(·)` (bound extraction)
  - `lower_inc(·)` (bound closed?), `lower_inf(·)` (unbounded?)
  - `isempty(·)`



## 10 | Geometric Objects

### Constructing **geometric objects** in PostgreSQL:



- Alternative string literal syntax (see PostgreSQL docs):
  - `'((x1,y1),(x2,y2))'::lseg`, `'<(x,y),r>'::circle`, ...

## Querying Geometric Objects

---


- A vast library of geometric operations (excerpt):

	Operation			Operation
<code>+, -</code>	translate		<code>area(.)</code>	area
<code>*</code>	scale/rotate		<code>height(.)</code>	height of box
<code>@-@</code>	length/circumference		<code>width(.)</code>	width of box
<code>@@</code>	center		<code>bound_box(.,.)</code>	bounding box
<code>&lt;-&gt;</code>	distance between		<code>diameter(.)</code>	diameter of circle
<code>&amp;&amp;</code>	overlaps?		<code>center(.)</code>	center
<code>&lt;&lt;</code>	strictly left of?		<code>isclosed(.)</code>	path closed?
<code>?- </code>	is perpendicular?		<code>npoints(.)</code>	# of points in path
<code>@&gt;</code>	contains?		<code>pclose(.)</code>	close an open path

- `<p>[0]`, `<p>[1]` to access x/y coordinate of point *p*.

## 🔧 Use Case: Shape Scanner

---



- Given an unknown shape (a **polygon** geometric object):
  - Perform horizontal “scan” to trace minimum/maximum (*i.e.*, bottom/top)  $y$  values for each  $x$ .
  - Use bottom/top traces to render the shape.

## 11 | JSON (JavaScript Object Notation)

---

**JSON** defines a textual data interchange format. Easy for humans to write and machines to parse (see <http://json.org>):

```

<object>    ::= { } | { <members> }
<members>   ::= <pair> | <pair> , <members>
<pair>      ::= <string> : <value>
<array>     ::= [ ] | [ <elements> ]
<elements>  ::= <value> | <value> , <elements>
<value>     ::= <string> | <number> | true | false | null
              | <array> | <object>

```

- SQL:2016 defines SQL↔JSON interoperability. JSON *<value>*s may be constructed/traversed and held in table cells (we still consider 1NF to be intact).

## JSON Sample *<value>s*

---

*<members>*

{ "title": "The Last Jedi", "episode": 8 }

↑  
*<object>*
↑  
*<pair>*

Table T (see Chapter 02):

*<elements>* { [ { "a": 1, "b": "x", "c": true, "d": 10 },  
                   { "a": 2, "b": "y", "c": true, "d": 40 },  
                   { "a": 3, "b": "x", "c": false, "d": 30 },  
                   { "a": 4, "b": "y", "c": false, "d": 20 },  
                   { "a": 5, "b": "x", "c": true, "d": null } ] }

↑  
*<number>*
↑  
*<array>* (of *<object>s*)


## JSON in PostgreSQL: Type `jsonb`<sup>3</sup>

---

Literal string syntax embeds JSON *<value>*s in SQL queries. Casting to type `jsonb` validates and encodes JSON syntax:

```
VALUES (1, '{ "b":1, "a":2 }' ::jsonb),  
       (2, '{ "a":1, "b":2, "a":3 }' ),  
       (3, '[ 0, false,null ]' );
```

column1	column2
1	{"a": 2, "b": 1}
2	{"a": 3, "b": 2}
3	[0, false, null]

<sup>3</sup> Alternative type `json` preserves member order, duplicate fields, and whitespace.  
 Reparses JSON values on each access, no index support.

## Navigating JSON *<value>s*

---

- **Access** field *f* / element at index *i* in array *<value>* via *->* or *->>:*<sup>4</sup>

<i>&lt;value&gt;-&gt;&lt;f&gt;</i>	}	yields a jsonb value, permits further navigation steps via <i>-&gt;</i> , <i>-&gt;&gt;</i>
<i>&lt;value&gt;-&gt;&lt;i&gt;</i>		
<i>&lt;value&gt;-&gt;&gt;&lt;f&gt;</i>	}	yields a text value (cast to atomic type for further computation)
<i>&lt;value&gt;-&gt;&gt;&lt;i&gt;</i>		

- **Path navigation:** chain multiple navigation steps via *#>* or *#>>: <value> #> '{<f or i>, ..., <f or i>}'*.

<sup>4</sup> Extracting non-existing fields yields *NULL*. JSON arrays are 0-based.

## Bridging between JSON and SQL

---

Turn the **fields and/or nested values** inside JSON object

$\langle o \rangle \equiv \{ \langle f_1 \rangle : \langle v_1 \rangle, \dots, \langle f_n \rangle : \langle v_n \rangle \}$  or array

$\langle a \rangle \equiv [\langle v_1 \rangle, \dots, \langle v_n \rangle]$  **into tables** which we can query:<sup>5</sup>

```
SELECT *
FROM  jsonb_each(<o>)
```

key	value
$\langle f_1 \rangle$	$\langle v_1 \rangle$
$\vdots$	$\vdots$
$\langle f_n \rangle$	$\langle v_n \rangle$

```
SELECT *
FROM  jsonb_array_elements(<a>)
```

value
$\langle v_1 \rangle$
$\vdots$
$\langle v_n \rangle$

<sup>5</sup> Re `jsonb_each(.)`: `jsonb_to_record(.)` or `jsonb_populate_record( $\tau$ ,.)` help to create typed records.



## Constructing JSON *<value>s*

---

- `row_to_json(·)::jsonb`

Convert a single **SQL row into a JSON *<object>***. Column names turn into JSON field names:

```
SELECT row_to_json(t)::jsonb -- yields objects of the form
FROM   T AS t;              -- {"a":·,"b":·,"c":·,"d":·}
```

- `array_to_json(array_agg(·))::jsonb`

Aggregate **JSON *<object>s* into a JSON *<array>***:

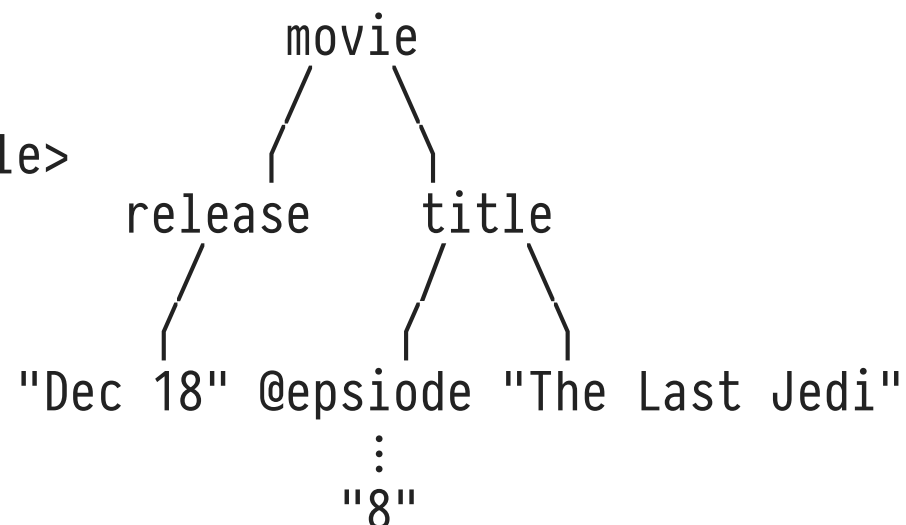
```
--   a unity for now, see Chapter 04
--   └──────────────────────────┘
SELECT array_to_json(array_agg(row_to_json(t))::jsonb
FROM   T AS t;
```

## 12 : XML (Extensible Markup Language)

---

**XML** defines a textual format describing ordered  $n$ -ary trees:

```
<movie>
  <release>Dec 18, 2017</release>
  <title episode="8">The Last Jedi</title>
</movie>
```



- XML support in SQL predates JSON support. Both are similar in nature. XML not discussed further here.<sup>6</sup>

<sup>6</sup> See the course [Database-Supported XML Processors](#).

## 13 | Sequences

---

**Sequences** represent counters of type **bigint** ( $-2^{63} \dots 2^{63} - 1$ ). Typically used to implement row identity/name generators:

```
CREATE SEQUENCE <seq>           -- sequence name
  [ INCREMENT <inc> ]           -- advance by <inc> (default: 1≡↑)
  [ MINVALUE <min> ]           -- range of valid counter values
  [ MAXVALUE <max> ]           -- (defaults: [1...263-1])
  [ START <start> ]             -- start (default: ↑<min>, ↓<max>)
  [ [NO] CYCLE ]                -- wrap around or error(≡ default)?
```

- Declaring a column of type **serial** creates a sequence:

```
CREATE TABLE <T> (... , <c> serial, ...) -- implies NOT NULL
      ↓
CREATE SEQUENCE <T>_<c>_seq;
```

## Advancing and Inspecting Sequence State

---

- Counter state can be (automatically) advanced and inspected:

```
CREATE SEQUENCE <seq> START 41 MAXVALUE 100 CYCLE;
:
SELECT nextval(' <seq> ');      -- ⇒ 41
SELECT nextval(' <seq> ');      -- ⇒ 42
SELECT currval(' <seq> ');      -- ⇒ 42
SELECT setval (' <seq> ',100);  -- ⇒ 100 (+ side effect)
SELECT nextval(' <seq> ');      -- ⇒ 1   (wrap-around)
```



sequence/table names are not 1<sup>st</sup> class in SQL

- Columns of type **serial** automatically populate with (and advance) their current counter value when set to **DEFAULT**.