

# Advanced SQL

---

04 — Arrays and User-Defined Functions

Summer 2020

Torsten Grust  
Universität Tübingen, Germany

## 1 | Arrays: Aliens(?) Inside Table Cells

---

SQL tables adhere to the **First Normal Form** (1NF): values  $v$  inside table cells are *atomic* w.r.t. the tabular data model:

...	A	...
...	$v$	...

Let us now discuss the **array** data type:

- $v$  may hold an ordered array of elements  $\{x_1, \dots, x_n\}$ .<sup>1</sup>
- SQL treats  $v$  as an atomic unit, but ...
- ... array functions and operators also enable SQL to query the  $x_i$  individually (still, that's no  $\nrightarrow$  with 1NF).

<sup>1</sup> To the PostgreSQL developer who decided to use  $\{\dots\}$  to denote *arrays*: **No dessert for you today!**

## 2 | Array Types

---

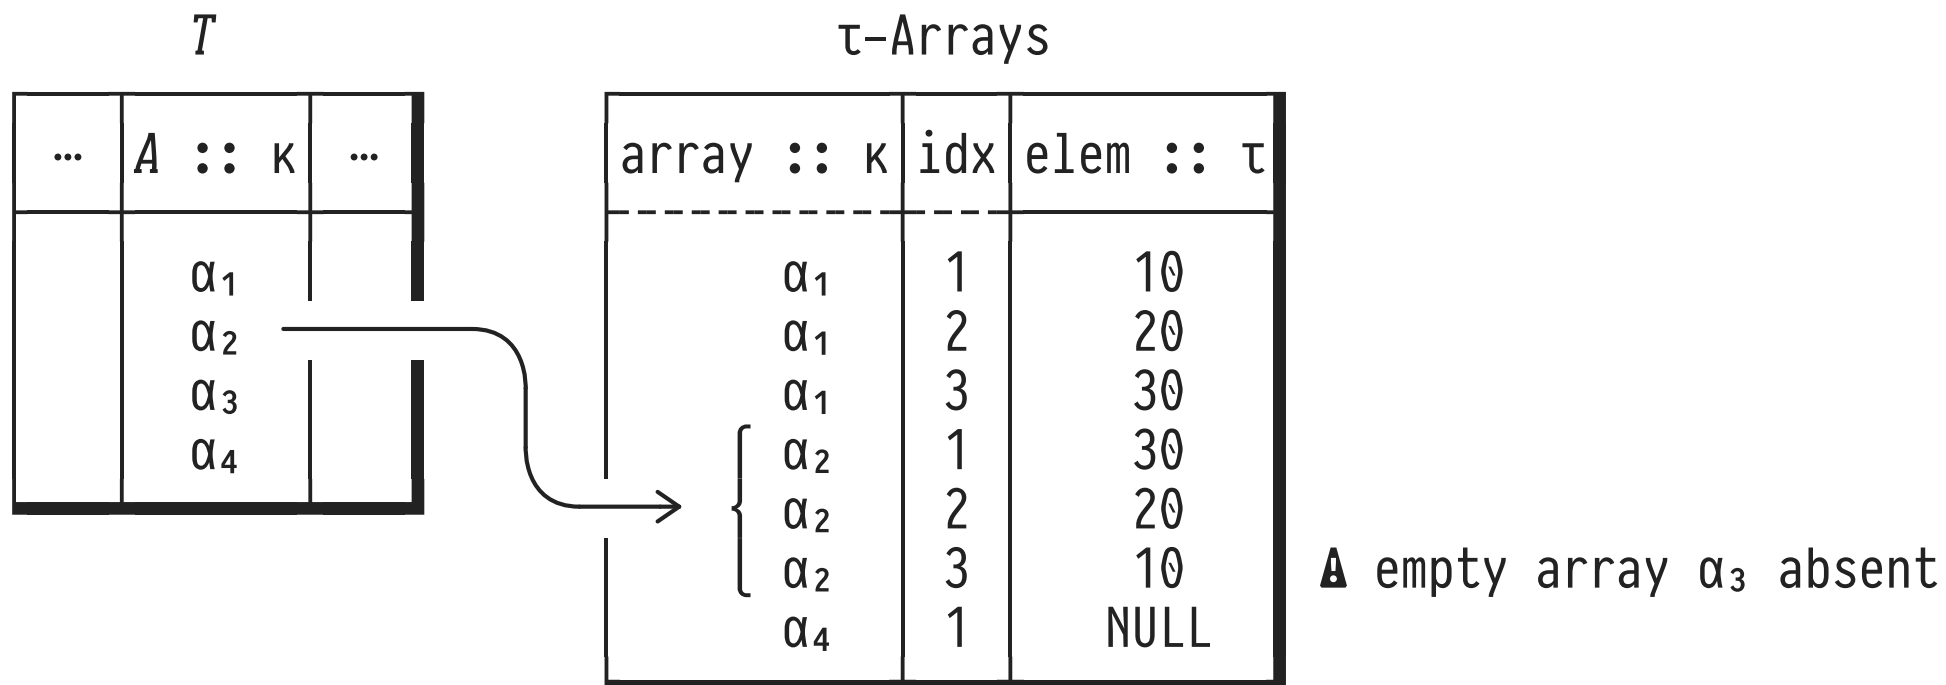
- For type  $\tau$ ,  $\tau[]$  (or  $\tau$  array) is the type of **homogenous arrays of elements of  $\tau$** .
  - $\tau$  may be built-in or user-defined (enums, row types).
  - Array size is unspecified—the array is dynamic.  
(PostgreSQL accepts  $\tau[n]$  but the  $n$  is ignored.)

...	$A :: \text{int}[]$	...
...	$\{10, 20, 30\}$	...
...	$\{30, 20, 10\}$	...
...	$\{\}$	...
...	$\{\text{NULL}\}$	...

$T$

# “Simulating” Arrays (Tabular Array Semantics)

---



- $\kappa$  denotes a suitable key data type.
- Arrays indexes are of type `int` and 1-based.

### 3 | Array Literals


---

#### One-dimensional array literals of type $\tau[]$ :

`array[] ::  $\tau$ []`      empty array of elements of type  $\tau$

`array[ $\langle x_1 \rangle, \dots, \langle x_n \rangle$ ]`       $\left. \begin{array}{l} \text{all } x_i \text{ of type } \tau \\ \text{'}\{\langle x_1 \rangle, \dots, \langle x_n \rangle\}\text{' :: } \tau[] \end{array} \right\}$

#### Multi-dimensional rectangular array literals of type $\tau[][]$ :

 all sub-arrays need to agree in size


`array[array[ $\langle x_{11} \rangle, \dots, \langle x_{1n} \rangle$ ], ..., array[ $\langle x_{k1} \rangle, \dots, \langle x_{kn} \rangle$ ]]`

`'{\{\langle x_{11} \rangle, \dots, \langle x_{1n} \rangle\}, \dots, \{\langle x_{k1} \rangle, \dots, \langle x_{kn} \rangle\}}' :: \tau[][]`

1 ■■■■  
⋮ ■■■■  
k ■■■■  
1.....n

# Example: Tree Encoding ( $\text{parents}[i] \equiv \text{parent of node } i$ )

---



Tree shape and node labels held in separate in-sync arrays:

tree	parents	labels
$t_1$	{NULL, 1, 2, 2, 1, 5}	{'a', 'b', 'd', 'e', 'c', 'f'}
$t_2$	{4, 1, 1, 6, 4, NULL, 6}	{'d', 'f', 'a', 'b', 'e', 'g', 'c'}
$t_3$	{NULL, 1, NULL, 1, 3}	{'a', 'b', 'd', 'c', 'e'}
	1 2 3 4 5	1 2 3 4 5 ← @idx

Trees

## Constructing Arrays

---

- **Append/prepend** element  $\ast$  to array or
- **concatenate** arrays:

```
array_append (array[x1, ..., xn],  $\ast$ )   $\equiv$  array[x1, ..., xn,  $\ast$ ]  
array_prepend(array[x1, ..., xn],  $\ast$ )   $\equiv$  array[ $\ast$ , x1, ..., xn]
```

```
array_cat(array[x1, ..., xn],  
          array[y1, ..., ym])            $\equiv$  array[x1, ..., xn, y1, ..., ym]
```

- Overloaded operator **||** embraces all of the above:

```
xs ||  $\ast$   $\equiv$  array_append(xs,  $\ast$ )  
 $\ast$  || xs  $\equiv$  array_prepend(xs,  $\ast$ )  
xs || ys  $\equiv$  array_cat(xs, ys)
```

## Accessing Arrays: Indexing / Slicing

---

- Array **indexes**  $i$  are 1-based (let  $xs \equiv \text{array}[x_1, \dots, x_n]$ ):

$xs[i]$	$\equiv x_i$	$i \notin \{1, \dots, n\}: \text{NULL}$
$(\text{NULL})[i]$	$\equiv \text{NULL}$	
$xs[\text{NULL}]$	$\equiv \text{NULL}$	
$xs[i:j]$	$\equiv \text{array}[x_i, \dots, x_j]$	$i > j: \text{array}[]$
$xs[i: ]$	$\equiv \text{array}[x_i, \dots, x_n]$	
$xs[ :j]$	$\equiv \text{array}[x_1, \dots, x_j]$	

- Access **last element**  $x_n$ :

$xs[\text{array\_length}(xs, 1)]$	# of elements in dimension <b>1</b> : $n$
$xs[\text{cardinality}(xs)]$	
↑	
$\Sigma$ (# of elements) in all dimensions	



## Searching for Elements in Arrays

---

Indexing accesses array by position. **Searching** accesses arrays by **contents**, instead.

- Let  $xs \equiv \text{array}[x_1, \dots, x_{i-1}, *, x_{i+1}, \dots, x_{j-1}, *, x_{j+1}, \dots, x_n]$  and comparison operator  $\theta \in \{=, <, >, <>, <=, >=\}$ :

$$x \theta \text{ ANY}(xs) \equiv \exists i \in \{1, \dots, n\}: x \theta xs[i]$$

$$x \theta \text{ ALL}(xs) \equiv \forall i \in \{1, \dots, n\}: x \theta xs[i]$$

$$\text{array\_position}(xs, *) \equiv i \quad \text{if } * \text{ not found: NULL}$$

$$\text{array\_positions}(xs, *) \equiv \text{array}[i, j] \quad \text{if } * \text{ not found: array}[]$$

$$\text{array\_replace}(xs, *, \#) \equiv \text{array}[x_1, \dots, \underset{i}{\#}, \dots, \underset{j}{\#}, \dots, x_n]$$

## 4 | A Bridge Between Arrays and Tables: `unnest` & `array_agg`

```
SELECT t.elem
FROM   unnest(array[x1,...,xn]) AS t(elem)
```

$\equiv XS$

$\equiv$

elem
x <sub>1</sub>
⋮
x <sub>n</sub>

Table t

```
SELECT array_agg(t.elem) AS xs
FROM   (VALUES (x1),
              ⋮
              (xn)) AS t(elem)
```

$\equiv$

xs
{x <sub>1</sub> , ..., x <sub>n</sub> }

- `unnest(•)`: a *set-returning function*. More on that soon.
- ⚠ Preservation of order of the  $x_i$  is *not* guaranteed...

## Representing Order (Indices) As First-Class Values

```
SELECT t.*
FROM   unnest(array[x1,...,xn])
      WITH ORDINALITY AS t(elem,idx)
```

≡

elem	idx
x <sub>1</sub>	1
⋮	⋮
x <sub>n</sub>	n

recall ordered aggregates

```
SELECT array_agg(t.elem ORDER BY t.idx) AS xs
FROM   (VALUES (x1,1),
              ⋮
              (xn,n)) AS t(elem,idx)
```

≡

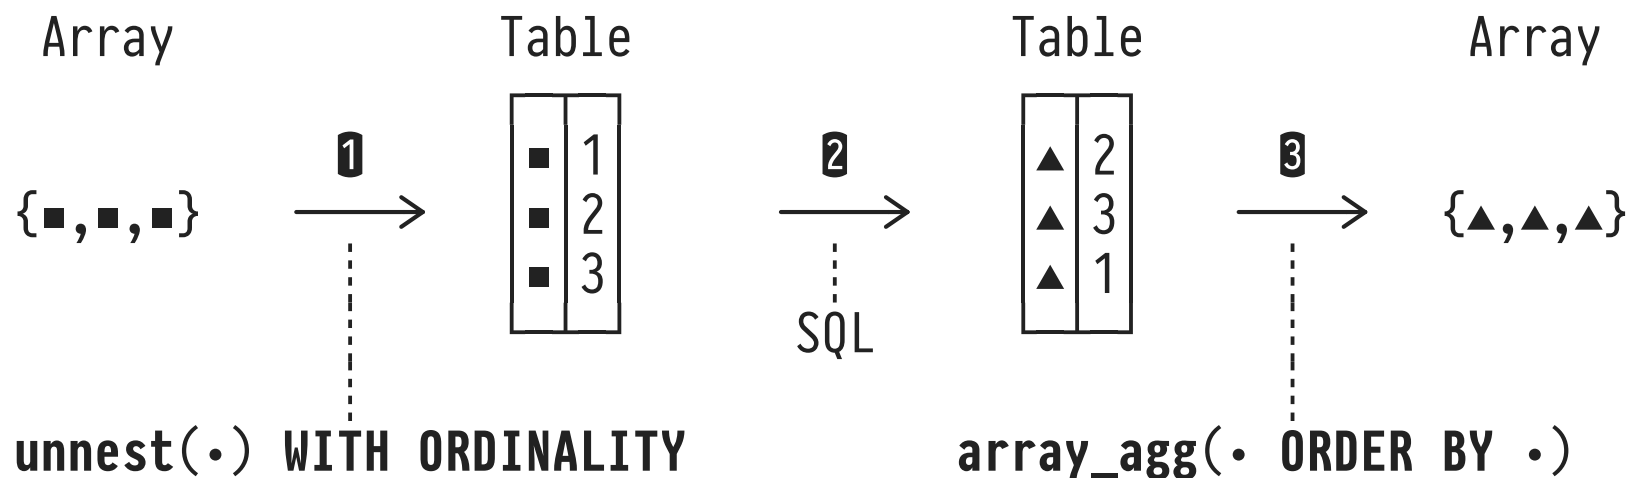
xs
{x <sub>1</sub> ,...,x <sub>n</sub> }

- `<f>(…)` `WITH ORDINALITY` adds a trailing column (▲) of ascending indices `1,2,...` to the output of function `f`.

## A Relational Array Programming Pattern

---

Availability of `unnest(·)` and ordered `array_agg(·)` suggests a pattern for **relational array programming**:



- At ❷ use the full force of SQL, read/transform/generate elements and their positions at will.
- ❶+❸ constitute **overhead**: an RDBMS is *not* an array PL.

## 5 | Table-Generating Functions

---

What is the **type** of `unnest(·)`?

- `unnest(·)` establishes a bridge between arrays and SQL's tabular data model:<sup>2</sup>

$\text{unnest} :: \tau[] \rightarrow \text{SETOF } \tau$
--

- In SQL, functions of type  $\tau_1 \rightarrow \text{SETOF } \tau_2$  are known as **set-returning** or **table(-generating) functions**. May be invoked wherever a query expects a table (`FROM` clause).
- Several built-in, but may also be **defined by the user**.

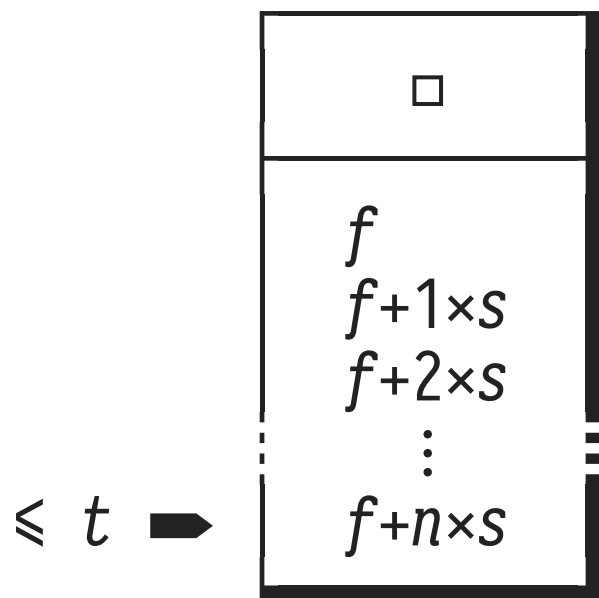
<sup>2</sup> Unfortunate naming again: `SETOF` should probably read `BAGOF` or `TABLE OF`.

## Series and Subscript Generators

---

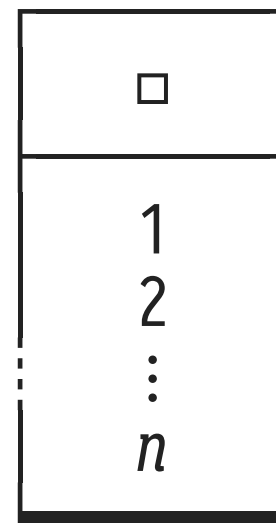
Built-in table-generating functions that generate **tables of consecutive numbers**:

**generate\_series(<f>, <t>, <s>)**



$\langle s \rangle \equiv 1$ , if absent  
 $\langle f \rangle, \langle t \rangle$ : numbers/timestamps

**generate\_subscripts(<xs>, <d>)**



$n \equiv \text{array\_length}(\langle xs \rangle, \langle d \rangle)$   
 can also enumerate  $n, \dots, 1$

## Text Generators (Regular Expression Matching)

---

Use *regular expression*<sup>3</sup> *re* to extract **matched substrings** from *t* or **split** text *t* at defined positions:

1. `regexp_matches(<t>,<re>,'g')`, yields SETOF `text[]`:  
Generates one array *xs* per match of *re* in *t*. Element `<xs>[i]` holds the **match** of the  $i^{\text{th}}$  *capture group* (in (...)).
2. `regexp_split_to_table(<t>,<re>)`, yields SETOF `text`:  
Uses the matches of *re* in *t* as *separators* to **split** *t*.  
Yields table of  $n+1$  rows if *re* matches  $n$  times.

<sup>3</sup> See [regexr.com](http://regexr.com) for tutorials and an interactive playground, for example.

## Breaking Bad: Parse a Chemical Formula (e.g., $\text{C}_6\text{H}_5\text{O}_7^{3-}$ )

---

```

SELECT t.match[1] AS element,  -- } extract match details
      t.match[2] AS "# atoms",  -- } from the (...)
      t.match[3] AS charge      -- } (capture groups)
FROM   regexp_matches(
      'C6H5O73-',
      '([A-Za-z]+)([0-9]*)([0-9]+[+-])?',
      'g')
      AS t(match);

```

element	# atoms	charge
C	6	NULL
H	5	NULL
O	7	3-

} NULL  $\equiv$  no match

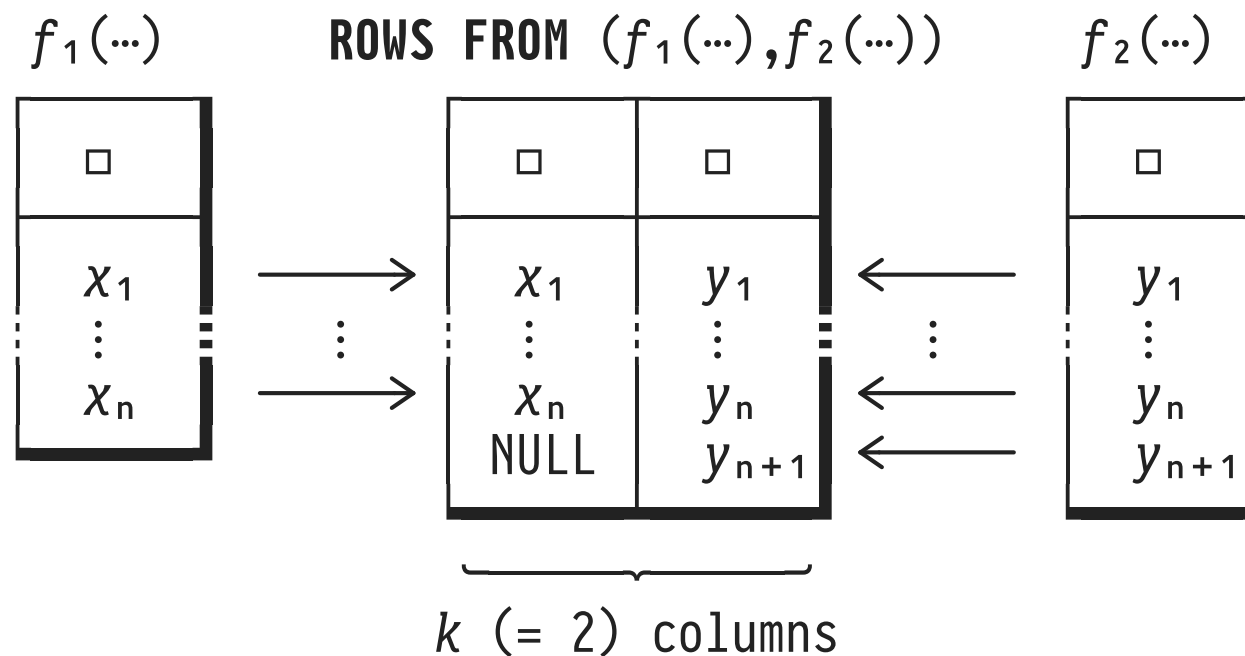


## Zippping Arrays and Table-Generating Functions

---

**Zip:** pair elements based on position (“ORDINALITY join”):

- Zippping table functions  $f_i$ : **ROWS FROM**( $f_1(\dots), \dots, f_k(\dots)$ )
- Zippping arrays  $xs_i$ : **unnest**( $xs_1, xs_2, \dots, xs_k$ )



## 6 : User-Defined SQL Functions (UDFs)

---

The body of a **user-defined SQL function (UDFs)** evaluates  $n \geq 1$  arbitrary SQL statements in sequence:

```
CREATE FUNCTION  $\langle f \rangle$ ( $\langle X_1 \rangle \tau_1, \dots, \langle X_k \rangle \tau_k$ ) RETURNS  $\tau$  AS
$$
   $\langle q_1 \rangle$ ;      --
   $\langle q_2 \rangle$ ;      -- } evaluate the  $q_i$  in order,
  :              -- :  $q_n$  defines the result
   $\langle q_n \rangle$       -- ]
$$
LANGUAGE SQL [IMMUTABLE];
```


↑

all  $q_i$  are read-only  $\Rightarrow f$  is free of side effects

- UDF  $f$  is stored persistently. Remove via **DROP FUNCTION**.

## UDF Types

---

- UDF  $f$  is  $k$ -ary with type  $\tau_1 \times \dots \times \tau_k \rightarrow \tau$ .
  - **Argument types**  $\tau_i$  must be **atomic** or **row types**.
  - **Overloading** allowed as long as  $(f, \tau_1, \dots, \tau_k)$  is unique.
  - Limited form of **polymorphism**: any  $\tau_i$  and  $\tau$  may be `anyelement/anyarray/anyenum/anyrange`.
    -  If `any...` occurs more than once in the function signature, *all* occurrences denote the *same* type:

$$\begin{array}{l}
 f_1 :: \overbrace{\text{anyelement} \times \text{anyelement}}^{\text{elem}} \rightarrow \text{boolean} \\
 f_2 :: \text{anyarray} \times \underbrace{\text{integer}}_{\text{elem}} \rightarrow \text{anyelement}
 \end{array}$$

## UDFs Can Return Tables

---

A UDF  $f :: \tau_1 \times \dots \times \tau_k \rightarrow \tau$  may be of **two flavors**:

	atomic $\tau$	$\tau \equiv \text{SETOF } \tau'$
If $q_n$ <sup>4</sup> returns no rows, If $q_n$ returns rows, May be invoked	returns <b>NULL</b> returns the first row wherever $v::\tau$ is used	returns empty table returns all rows in the <b>FROM</b> clause

### Regular vs. Table-generating UDFs

- A UDF may invoke **INSERT/DELETE/UPDATE** statements in  $q_i$  and thus incur side-effects. (Hmm, UDF...⚡)
  - No **IMMUTABLE** option—use **VOLATILE** instead.
  - Use  $\tau \equiv \text{void}$  if  $f$  is all about side-effects or consider adding **... RETURNING  $\langle e_1 \rangle, \dots, \langle e_m \rangle$**  if  $i = n$ .

<sup>4</sup> Recall:  $f$ 's body evaluates queries  $q_1, \dots, q_n$  (in this order).

## Example UDF: Map Unicode Subscripts

---

Map subscript symbol `'0'`, ..., `'9'` to its value in `{0,...,9}`:

```
CREATE FUNCTION subscript(s text) RETURNS int AS
$$
  SELECT subs.value::int - 1
  FROM   unnest(array['0', '1', '2', ..., '9'])
        WITH ORDINALITY AS subs(sym, value)
  WHERE  subs.sym = s
$$
LANGUAGE SQL IMMUTABLE;
```

- This is a UDF with atomic return type: yields `NULL` if `s` does not denote a valid subscript.

## Example UDF: Issue Unique ID, Write Protocol

---

Generate ID of the form '<prefix>###' and log time of issue:

```
CREATE FUNCTION new_ID(prefix text) RETURNS text AS
$$
  INSERT INTO issue(id,"when") VALUES
    (DEFAULT, 'now'::timestamp)
  RETURNING prefix || id::text           -- id: just generated
$$
LANGUAGE SQL VOLATILE;                  -- function is side-effecting
```

id	when
⋮	⋮
42	2020-05-12 17:26:14.188803
⋮	⋮

Table `issue`

## Example Table-Generating UDF: Flatten a 2D-Array

---

Unnest 2D array `xss` in *column-major order*:<sup>5</sup>

```
CREATE OR REPLACE FUNCTION unnest2(xss anyarray)
  RETURNS SETOF anyelement AS
$$
  SELECT xss[i][j]
  FROM   generate_subscripts(xss,1) _(i),
         generate_subscripts(xss,2) __(j)
  ORDER BY j, i -- return elements in column-major order
$$
LANGUAGE SQL IMMUTABLE;
```

- ⚠ Intended type is `unnest2 ::  $\tau[][] \rightarrow \text{SETOF } \tau$` .

<sup>5</sup> Built-in function `unnest(.)` can flatten  $n$ -dimensional arrays in row-major order.

## Table-Generating UDFs: Returning Typed Rows

---

Assume a table-generating UDF  $f :: \dots \rightarrow \tau$ .

If  $\tau \equiv$

SETOF  $\tau'$   
 $\tau'$  atomic

$\square$
$v_1$
$\vdots$
$v_n$

$v_i :: \tau'$

SETOF  $\tau'$   
 $\tau' \equiv (c_1 :: \tau_1, \dots, c_m :: \tau_m)$

$c_1$	$\dots$	$c_m$

TABLE ( $c_1 \ \tau_1, \dots, c_m \ \tau_m$ )

$c_1$	$\dots$	$c_m$

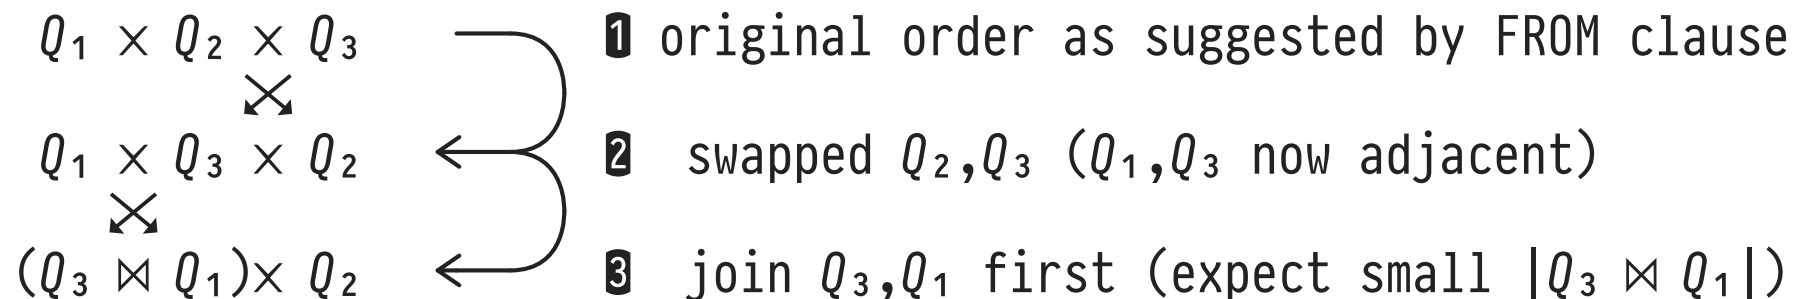
equivalent, but do not need named row type  $\tau'$



## 7 : ',' in the **FROM** Clause and Row Variable References

```
SELECT ...
FROM   Q1 AS t1, Q2 AS t2, Q3 AS t3 -- ti<j not free in Qj
```

- Q: Why is  $t_{i<j}$  not usable in  $Q_j$ ?
- A: “... the **','** in **FROM** is commutative and associative...”.  
Query optimization might rearrange the  $Q_j$ :



## But Dependent Iteration in **FROM** is Useful...

---

Recall (find largest label in each tree  $t_1$ ):

```

SELECT t1.tree, MAX(t2.label) AS "largest label"
--           Q1                Q2
--           └───┬──────────┘
FROM   Trees AS t1, unnest(t1.labels) AS t2(label)
GROUP BY t1.tree;
           ▲
           ⚡

```

- **Dependent iteration** (here:  $Q_2$  depends on  $t_1$  defined in  $Q_1$ ) has its uses and admits intuitive query formulation.
- $\Rightarrow$  Exception: the arguments of table-generating functions may refer to row variables defined earlier (like  $t_1$ ).

## LATERAL:<sup>6</sup> Dependent Iteration for Everyone

---

Prefix  $Q_j$  with **LATERAL** in the **FROM** clause to announce dependent iteration:

```
SELECT ...
FROM    $Q_1$  AS  $t_1$ , ..., LATERAL  $Q_j$  AS  $t_j$ , ...
                        ▲
                        may refer to  $t_1, \dots, t_{j-1}$ 
```

- Works for *any* table-valued SQL expression  $Q_j$ , subqueries in (...) in particular.
  - Good style: be explicit and use **LATERAL** even with table functions.

<sup>6</sup> Lateral /'læt(ə)rəl/ a. [Latin lateralis]: *sideways*

## LATERAL: SQL's for each-Loop

---

LATERAL admits the formulation of **nested-loops** computation:

```
SELECT e
FROM   Q1 AS t1, LATERAL Q2 AS t2, LATERAL Q3 AS t3
```

is evaluated just like this nested loop:

```
for t1 in Q1
  for t2 in Q2(t1)
    for t3 in Q3(t1, t2)
      return e(t1, t2, t3)
```

- Convenient, intuitive, and perfectly OK.  
But much like hand-cuffs for the query optimizer. ⚠

## LATERAL Example: Find the Top $n$ Rows Among a Peer Group

---

Which are **the three tallest** two- and four-legged dinosaurs?

```
SELECT locomotion.legs, tallest.species, tallest.height
FROM   (VALUES (2), (4)) AS locomotion(legs),
       LATERAL (SELECT d.*
                FROM   dinosaurs AS d
                WHERE  d.legs = locomotion.legs ◀
                ORDER BY d.height DESC
                LIMIT 3) AS tallest
```

legs	species	height
2	Tyrannosaurus	7
2	Ceratosaurus	4
2	Spinosaurus	2.4
4	Supersaurus	10
4	Brachiosaurus	7.6
4	Diplodocus	3.6

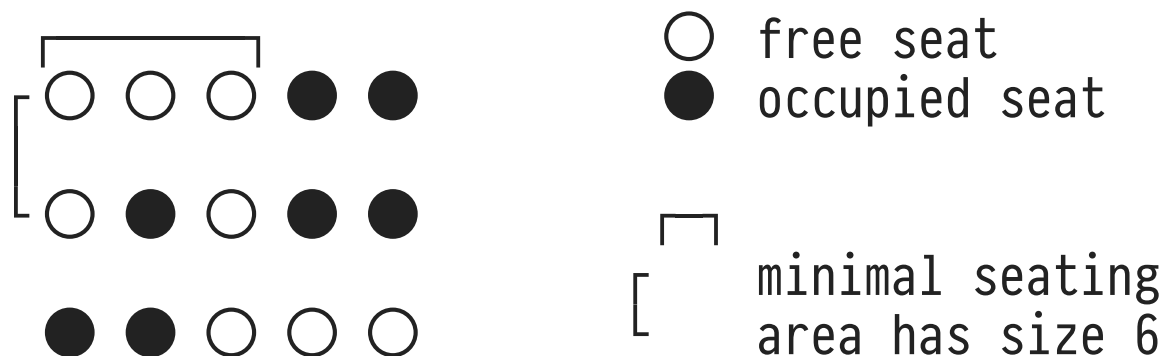
## 8 | ACM ICPC: Finding Seats

---

ACM ICPC Task **Finding Seats** (South American Regionals, 2007)

*“ $K$  friends go to the movies but they are late for tickets. To sit all nearby, they are looking for  $K$  free seats such that the rectangle containing these seats has minimal area.”*

- Assume  $K = 5$ :



## 🔧 Finding Seats: Parse the ICPC Input Format

---

- Typical ICPC character-based input format:

...XX <sup>C<sub>R</sub></sup>	. free seat
.X.XX <sup>C<sub>R</sub></sup>	X occupied seat
XX...	<sup>C<sub>R</sub></sup> new line

- **Parse into table** making seat position/status explicit:

<u>row</u>	<u>col</u>	<u>taken?</u>
1	1	false
1	2	false
1	3	false
1	4	true
⋮	⋮	⋮
3	5	false

Table `seats`

## Finding Seats: Parse the ICPC Input Format (Table **seats**)

---

```
\set cinema '...XX\\n.X.XX\\nXX...'

SELECT  row.pos, col.pos, col.x = 'X' AS "taken?"
FROM    -- rows
        unnest(string_to_array(:'cinema', '\n'))
        WITH ORDINALITY AS row(xs, pos),
        -- columns
        LATERAL unnest(string_to_array(row.xs, NULL))
                WITH ORDINALITY AS col(x, pos)
```

- `string_to_array(:'cinema', '\n')` yields an array of three row strings: `{'...XX', '.X.XX', 'XX...'}`.
- `string_to_array(row.xs, NULL)` splits string `row.xs` into an array of individual characters (= seats).



## 🔧 Finding Seats: A Problem Solution (Generate and Test)

---

- **Query Plan:**

1. Determine the extent (*rows* × *cols*) of the cinema seating plan.
2. **Generate all** possible north-west (*nw*) and south-east (*se*) corners of rectangular seating areas:
  - For each such 「*nw,se*」 rectangle, scan its seats and **test** whether the number of free seats is  $\geq K$ .
  - If so, record *nw* together with the rectangle's *width/height*.
3. Among these rectangles with sufficient seating space, select those with minimal area.

## 🔧 Finding Seats: Generating All Possible Rectangles

---

Generate all `「nw,se」` corners for rectangles up to maximum size `rows × cols`:

```
SELECT ROW(row_nw, col_nw) AS nw,
       ROW(row_se, col_se) AS se
FROM   generate_series(1, rows)           AS row_nw,
       generate_series(1, cols)          AS col_nw,
       LATERAL generate_series(row_nw, rows) AS row_se,
       LATERAL generate_series(col_nw, cols) AS col_se
```

Generates  $\begin{pmatrix} rows \\ \sum_{r=1} \end{pmatrix} \times \begin{pmatrix} cols \\ \sum_{c=1} \end{pmatrix}$  rectangles  $\Rightarrow$  test/filter early!