

Advanced SQL

05 — Window Functions

Summer 2020

Torsten Grust
Universität Tübingen, Germany

1 | Window Functions

With SQL:2003, the ISO SQL Standard introduced **window functions**, a new mode of row-based computation:


SQL Feature	Mode of Computation
function	row → row
table-generating function	row → table of rows
aggregate	group of rows → row (one per group)
window function 🍷	row vicinity → row (one per row)


SQL Modes of Computation

Window functions ...

- ... are **row-based**: each individual input row *r* is mapped to one result row,
- ... use the **vicinity** around *r* to compute this result row.

Row Vicinity: Window Frames

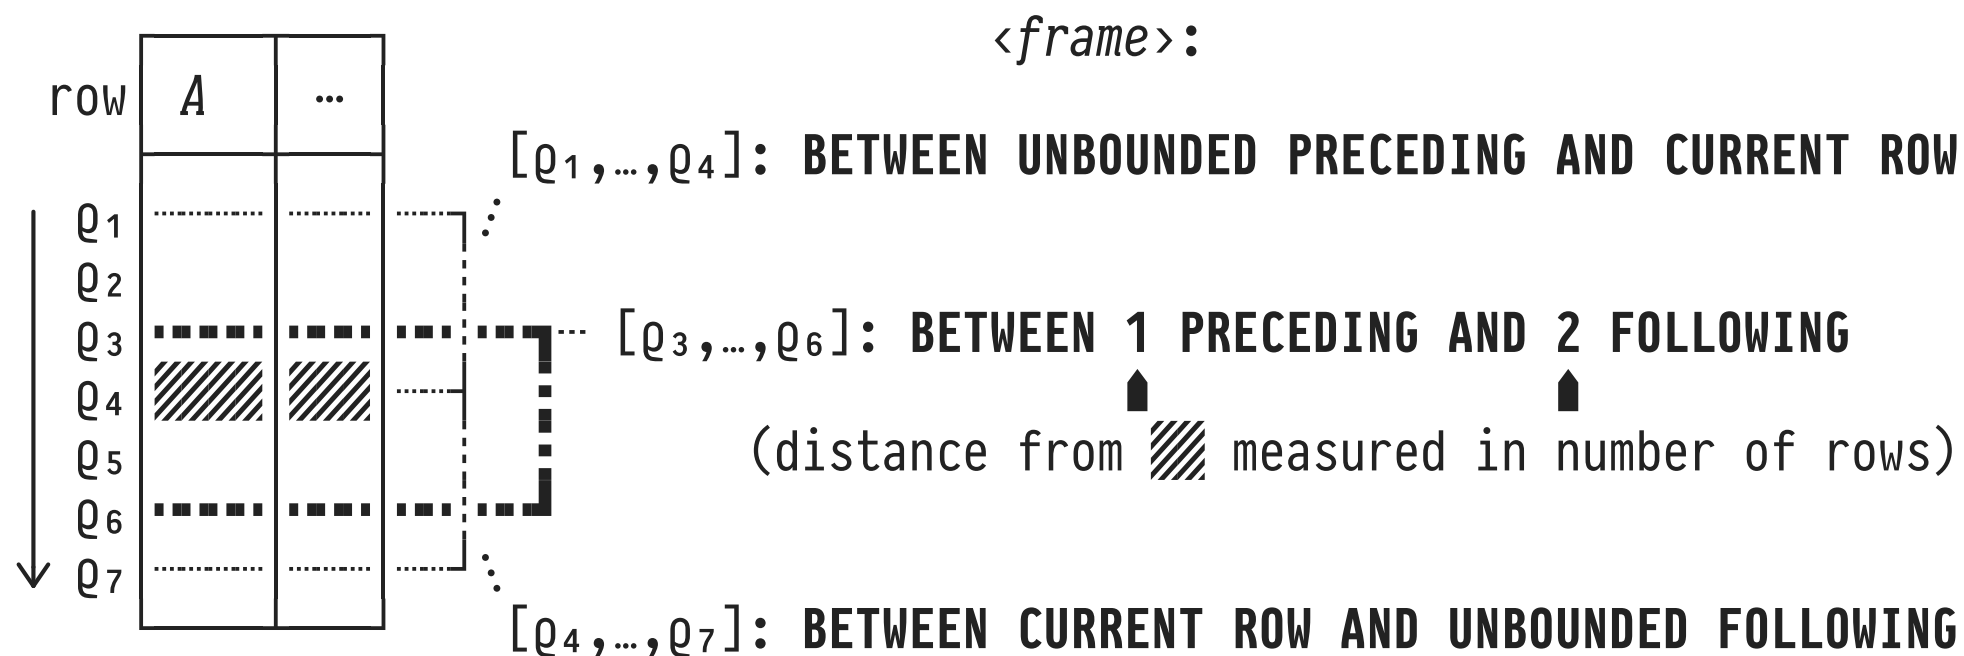


- Each row is the **current row**  at one point in time.
- Row vicinity (**window, frame**) is based on either:
 - ① row **position** (**ROWS** windows),
 - ② row **values** v_i (**RANGE** windows),
 - ③ row **peers** (**GROUPS** windows).

- As the current row changes, the window *slides* with it.
-  Window semantics depend on a defined **row ordering**.

Window Frame Specifications (Variant: **ROWS**)

window function ordering criteria frame specification
 $\underbrace{\hspace{1cm}}$ $\underbrace{\hspace{2cm}}$ $\underbrace{\hspace{1cm}}$
 $\langle f \rangle$ **OVER** (**ORDER BY** $\langle e_1 \rangle, \dots, \langle e_n \rangle$ [**ROWS** $\langle frame \rangle$])

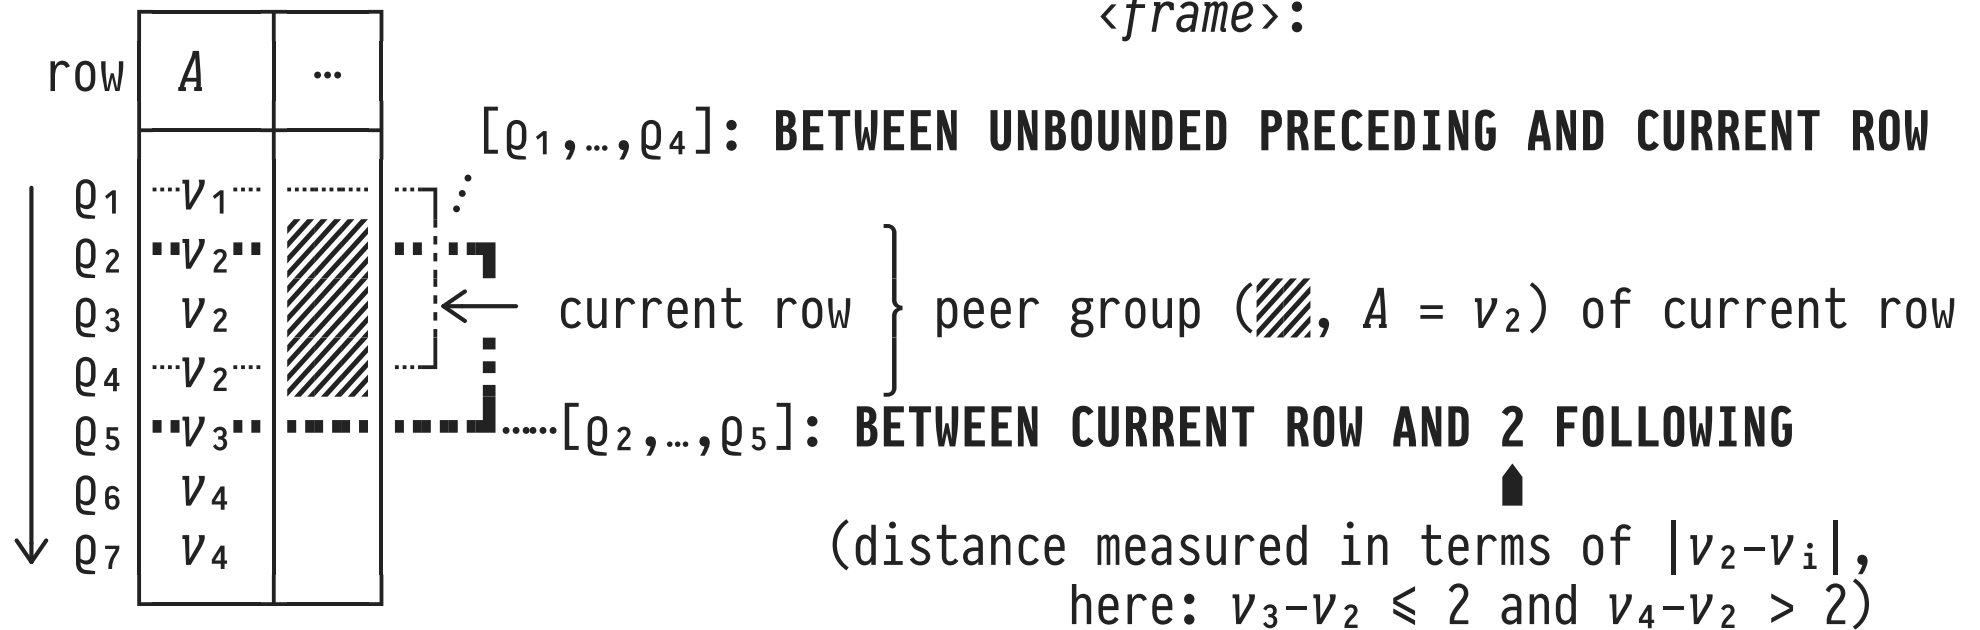


Window Frame Specifications (Variant: **RANGE**)

window function one column frame specification

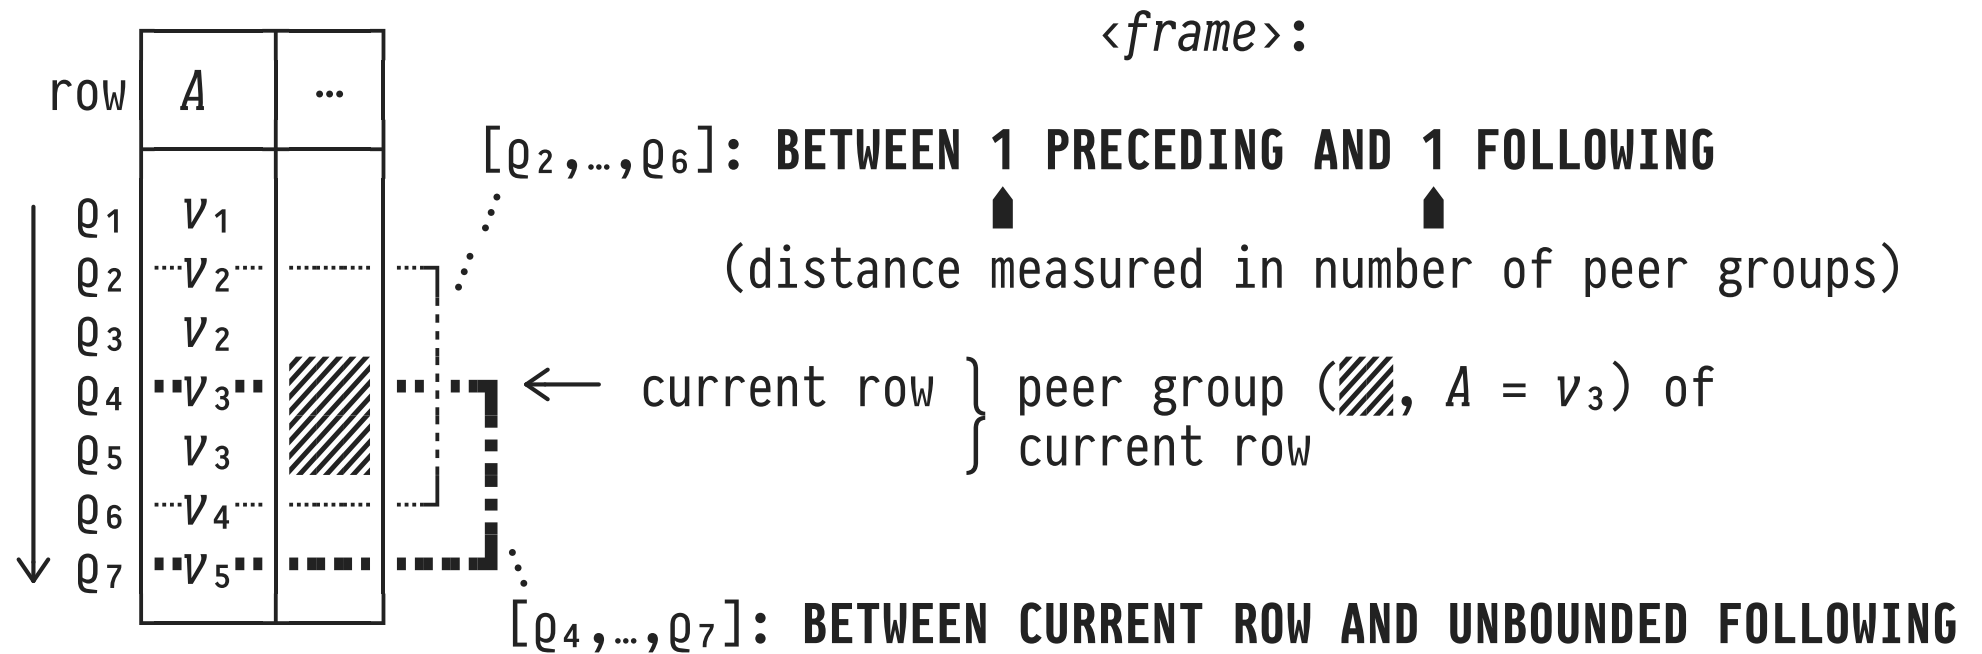
$\underbrace{\langle f \rangle}$ **OVER** (**ORDER BY** $\underbrace{\langle A \rangle}$ [**RANGE** $\underbrace{\langle frame \rangle}$])

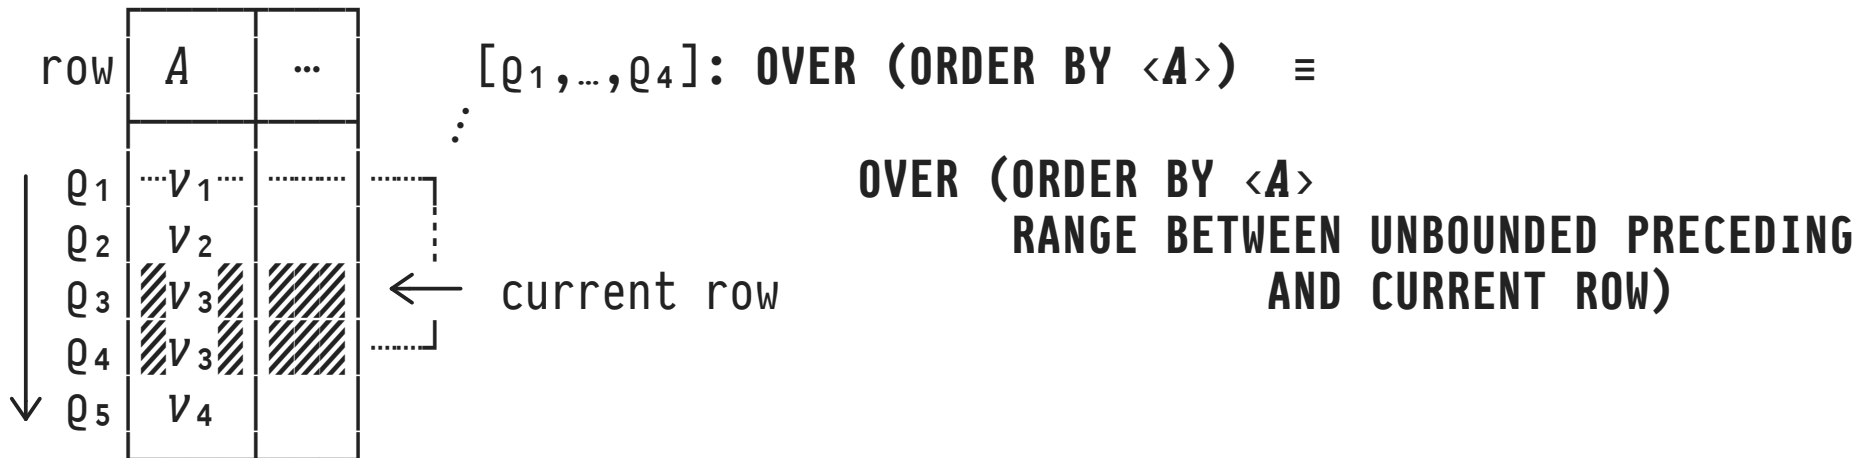
$\langle frame \rangle$:



Window Frame Specifications (Variant: **GROUPS**)

window function ordering criteria frame specification
 $\underbrace{\hspace{10em}}$ $\underbrace{\hspace{10em}}$ $\underbrace{\hspace{10em}}$
 $\langle f \rangle$ **OVER** (**ORDER BY** $\langle e_1 \rangle, \dots, \langle e_n \rangle$ [**GROUPS** $\langle frame \rangle$])





WINDOW Clause: Name the Frame

Syntactic Ⓢ: If window frame specifications

1. become unwieldy because of verbose SQL syntax and/or
2. one frame is used multiple times in a query,

add a **WINDOW** clause to a SFW block to **name the frame**, *e.g.*:

```
SELECT ... <f> OVER <wi> ... <g> OVER <wj> ...  
FROM ...  
WHERE ...  
⋮  
WINDOW <w1> AS (<frame1>), ..., <wn> AS (<framen>)  
ORDER BY ...
```


Use SQL Itself to Explain Window Frame Semantics

Regular **aggregates** may act as window functions $\langle f \rangle$. All **rows** in the frame will be aggregated:

```
SELECT w.row AS "current row",
       COUNT(*) OVER win AS "frame size",
       array_agg(w.row) OVER win AS "rows in frame"
FROM   W AS w
WINDOW win AS (<frame>)
```

<u>row</u>	a	b
q ₁	1	●
q ₂	2	○
q ₃	3	○
q ₄	3	●
⋮	⋮	⋮

Table W

🔧 Q: What is the Chance of Fine Weather on Weekends?

Input: Daily weather readings in `sensors`:

<u>day</u>	weekday	temp	rain
1	Fri	10	800
2	Sat	12	300
⋮	⋮	⋮	⋮

Table `sensors`

- The weather is fine on day *d* if—on *d* and the two days **prior**—the minimum temperature is above 15°C and the overall rainfall is less than 600ml/m².
- **Expected output:**

weekend?	% fine
f	29
t	43

2 | **PARTITION BY:** Window Frames Inside Partitions


Optionally, we may **partition** the input table *before* rows are sorted and window frames are determined:

all input rows that agree on all $\langle p_i \rangle$ form one partition

```
 $\langle f \rangle$  OVER ( [ PARTITION BY  $\langle p_1 \rangle, \dots, \langle p_m \rangle$  ]  
              [ ORDER BY  $\langle e_1 \rangle, \dots, \langle e_n \rangle$  ]  
              [  $\langle frame \rangle$  ] )
```

- Note:
 1. Frames **never cross partitions**.
 2. **BETWEEN ... PRECEDING AND ... FOLLOWING** respects **partition boundaries**.



🔧 Q: Which Spots are Visible in a Hilly Landscape?



- From the viewpoint of p_0 (🧑) we can see p_1 , but...
 - ... p_2 is **obscured** (no straight-line view from p_0),
 - ... p_3 is **obscured** (lies behind the 800m peak).

🔧 Q: Visible Spots in a Hilly Landscape? — A: MAX Scan!

$$\alpha = \text{atan}(\Delta\text{alt}/\Delta x)$$



- We have $0^\circ < \alpha_1 < \alpha_2 < \alpha_3$ and $\beta \geq \alpha_2$, $\gamma < \alpha_2$, $\delta < \alpha_3$.
- \uparrow \uparrow \uparrow
 p_1 visible $p_{2,3}$ obscured

🔧 Q: Visible Spots in a Hilly Landscape? — A: MAX Scan!

- **Input:** Location of p_0 (here: $x = 0$) and 1D-map of hills:

x	alt
0	200
10	200
\vdots	\vdots
120	500

Table `map`

- **Output:** Can p_0 see the point on the hilltop at x ?

x	visible?
0	true
10	true
\vdots	\vdots
120	false

Q: Visible Spots in a Hilly Landscape? — A: MAX Scan!

WITH

-- ❶ Angles α (in $^\circ$) between p_0 and the hilltop at x
angles(x , angle) **AS** (
 SELECT $m.x$,

degrees(**atan**(($m.alt - p_0.alt$) /
 abs($p_0.x - m.x$))) **AS** angle

FROM map **AS** m

WHERE $m.x > p_0.x$),

-- ❷ MAX(α) scan (to the right of p_0)

max_scan(x , max_angle) **AS** (
 SELECT $a.x$,


MAX($a.angle$)

OVER (**ORDER BY** $abs(p_0.x - a.x)$) **AS** max_angle

FROM angles **AS** a),

:

Looking Left *and* Right: PARTITION BY



- Need MAX scans left *and* right of $p_0 \Rightarrow$ use PARTITION BY.

Looking Left *and* Right: **PARTITION BY**

```

WITH
:
-- 2 MAX( $\alpha$ ) scan (left/right of  $p_0$ )
max_scan(x, max_angle) AS (
  SELECT a.x,                                --  $\in \{-1, 0, 1\}$ 
         MAX(a.angle)                        --  $\underbrace{\hspace{1.5cm}}$ 
         OVER (PARTITION BY sign( $p_0.x - a.x$ )
              ORDER BY abs( $p_0.x - a.x$ )) AS max_angle
  FROM   angles AS a  --  $\underbrace{\hspace{1.5cm}}$ 
                   --  $\Delta x > 0$ 
),
:

```

- $\forall a \in \text{angles}: a.x \neq p_0.x \Rightarrow$ We end up with **two** partitions.

3 | Scans: Not Only in the Hills

Scans are a general and expressive computational pattern:

$\underbrace{\langle agg \rangle(\langle e \rangle)}_{(\phi, z, \oplus)} \text{ OVER } (\text{ORDER BY } \langle e_1 \rangle, \dots, \langle e_n \rangle$	$\{\text{ROWS, RANGE, GROUPS}\} \text{ BETWEEN}$
	$\text{UNBOUNDED PRECEDING AND CURRENT ROW})$

- Available in a variety of forms in programming languages
 - Haskell: `scanl z ⊕ xs`, APL: `⊕\xs`, Python: `accumulate`:
`scanl ⊕ z [x1, x2, ...] = [z, z ⊕ x1, (z ⊕ x1) ⊕ x2, ...]`
- In parallel programming: *prefix sums* (👉 Guy Blelloch)
 - Sorting, lexical analysis, tree operations, reg.exp. search, drawing operations, image processing, ...

4 | Interlude: Quiz

Q: Assume $xs \equiv '((b*2)-4*a*c)*0.5'$. What is computed below?

```
SELECT inp.pos, inp.c,
       SUM((array[1,-1,0])[COALESCE(p.oc, 3)])
         OVER (ORDER BY inp.pos) AS d
FROM   unnest(string_to_array(xs, NULL))
       WITH ORDINALITY AS inp(c,pos),
       LATERAL (VALUES (array_position(array['(', ')'],
                                       inp.c))) AS p(oc)
ORDER BY inp.pos;
```

💡 **Hint** (this is the same query expressed in APL):

```
xs ← '((b*2)-4*a*c)*0.5'
+ \ (1 -1 0) ['(') ⌊ xs]
```