

Advanced SQL

07 — Procedural SQL

Summer 2020

Torsten Grust
Universität Tübingen, Germany

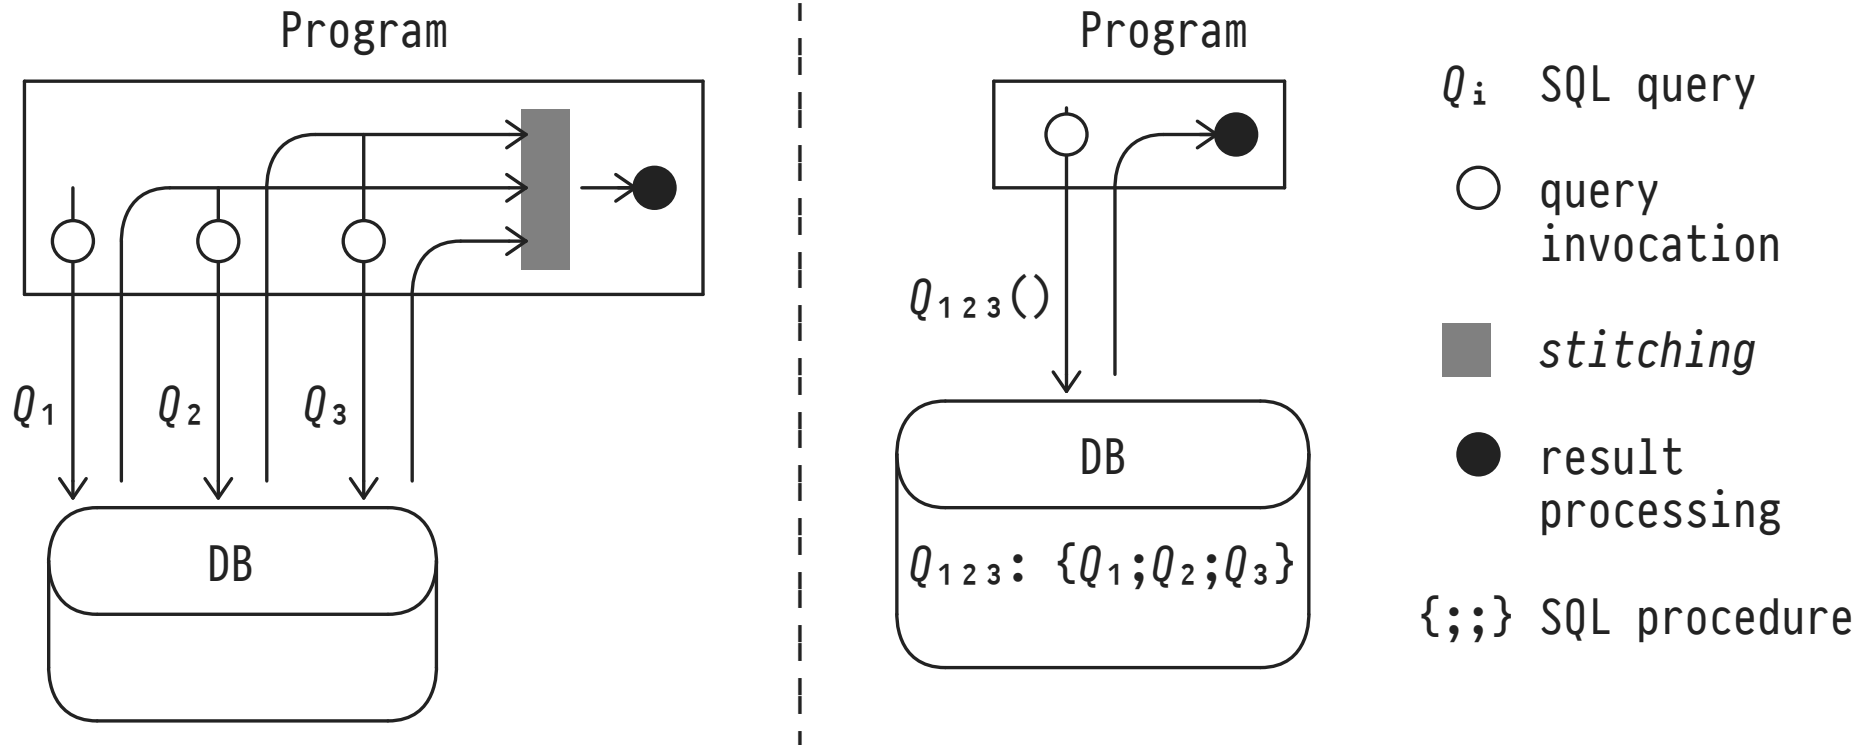
1 | Scripting Language + SQL = Procedural SQL

We started out in this course with the aim to **move more computation close to the data**. Admitting recursion in SQL is one way to declaratively express complex computation.

Procedural SQL follows an entirely different path towards this goal:

- Implement application logic *inside* the RDBMS, even if this **computation is inherently procedural** (= sequential, imperative).
- Use **SQL as a sub-language of a scripting language** whose types match those of the tabular data model.

Procedural SQL: Less Round-Trips, Less Stitching



- *Stitching*: On the PL heap, piece together the tabular results delivered by the individual SQL queries Q_i .

Procedural SQL aka *Stored Procedures*

Code in Procedural SQL is organized in **functions/procedures that are stored persistently** by the DBMS.¹

These functions/procedures...

- may be used anywhere that SQL's built-ins could be used,
- inherit all user-defined types, functions, and operators,
- can define new operators, aggregate/window functions, and triggers.

¹ This implies that we need to manage these procedures using familiar constructs like `CREATE PROCEDURE ...`, `CREATE FUNCTION ...`, `DROP PROCEDURE [IF EXISTS] ...`, etc.

PL/SQL:² Scripting with SQL Types

```
CREATE FUNCTION  $f(x_1 \ \tau_1, \dots, x_n \ \tau_n)$  RETURNS  $\tau$  AS
 $\$ \$ \dots \langle block \rangle \dots \$ \$$ 
LANGUAGE PLPGSQL;
```

- The τ_i , τ may be any scalar, array, or (named) row type.
- Limited polymorphism: functions may accept/return types `anyelement`, `anyarray` (recall our discussion of SQL UDFs).
 - Functions may return type `record` (then the caller must provide column names/types through explicit aliasing).
- Functions may return—but *not* accept \rightsquigarrow —sets of (row) values with $\tau \equiv \text{SETOF } \bar{\tau}$.

² *PL/SQL* is the widely adopted abbreviation for *Procedural Language for SQL*, originating in the Oracle® RDBMS. Variants include *Transact-SQL* (Microsoft® SQL Server) and *PL/pgSQL* (PostgreSQL).

2 | Block Structure

PL/SQL code is organized in (nested) **blocks** that group statements and define **variable scopes**:

```

block {
    [ DECLARE <declarations> ]
    BEGIN
        <statement>           -- ← any statement may be
    END;                     -- a (sub-)block again

```

- Declared variables are in scope in the block and its sub-blocks. Local names shadow outer names.
- Optionally introduce block with `<< <label> >>`: variable `v` may then also be referred to as `<label>.v`.
- Outermost block of body for `f` has implicit `<< <f> >>`.

Block Structure and Variable Scope³

<i>in scope</i>		CREATE FUNCTION $f(x_1 \ \tau_1)$ RETURNS τ AS	
		\$\$	
$f.x_1$	[<< 0 >>	-- outer block
$f.x_1, \ 0.v$		DECLARE $v \ \tau_v$;	
	[BEGIN	
		\vdots	
	[<< i >>	-- inner (sub-)block
		DECLARE $v \ \tau_u$;	
$f.x_1, \ 0.v, \ i.v$	[BEGIN	
		\vdots	
]	END;	
		\vdots	
]	END;	
		\$\$ LANGUAGE PLPGSQL;	

³ Additional special variables (like `FOUND`) are bound in the outermost f scope (see below).

3 | Variable Declarations

The optional `DECLARE <declarations>` brings **typed variable(s)** `v` into scope. An initial binding expression `e` may be given:

```
DECLARE v [ CONSTANT ] τ [ NOT NULL ] [ := e ];  
      ⋮
```

- If `:= e` is omitted, `v` has initial value `NULL`.
- `NOT NULL`: any assignment of `NULL` yields a runtime error.
- `CONSTANT`: the initial binding may not be overwritten.
- Use `c%TYPE` for `τ` to declare `v` with the same type as variable or table column named `v`.

Variables With Row Types Have Row Values

Let T be a table with **row type** $(c_1 \tau_1, \dots, c_n \tau_n)$. Recall: this row type is also known as T . Thus:

```
--                                ↓ row type name
CREATE FUNCTION accessi(t T) RETURNS T.ci%TYPE AS
$$
--          ↓+↓ table + column name
DECLARE x T.ci%TYPE; -- x has type  $\tau_i$ 
BEGIN
    x := t.ci;          -- field access uses dot notation
    RETURN x;
END;
$$
LANGUAGE PLPGSQL;
```

4 : PL/SQL Expressions

In PL/SQL, any expression e that could also occur in a **SELECT** clause, is a valid expression.

In fact, the execution of PL/pgSQL statements like

```
 $v := e$   
IF  $e$  THEN ... ELSE ... END IF
```

lead to the evaluation of **SELECT** e by the SQL interpreter.

- Interoperability between PL/pgSQL and SQL. 👍
- Performance impact: context switches PL/SQL↔SQL. 👎
- If $e \equiv e(x,y)$, compile SQL once with parameters x,y .

5 | PL/SQL Statements — Assignment

$v := e$

1. Evaluate e , yields a single value (scalar, row, array, user-defined, including `NULL`). e may **not be table-valued**.
2. Cast value to type τ of v .
 - SQL casting rules apply (may fail at runtime).
 - e may use textual literal syntax (e.g., for user-defined enumerations, JSON, or geometric objects).
3. Bind variable v to value.

Assignment of Single-Row Query Results

A single-row⁴ SQL query augmented with **INTO** is a valid PL/SQL assignment statement:

1 SELECT e_1, e_2, \dots, e_n INTO v FROM \dots		2 SELECT e_1, e_2, \dots, e_n INTO v_1, v_2, \dots, v_n FROM \dots
---	--	--

1. Evaluate SQL query, obtain a single row of n values.
 - ① Assign row value to row-typed variable v , or
 - ② assign value of e_i to v_i ($i \in \{1, \dots, n\}$).
2. Variable **FOUND** :: **boolean** indicates if a row was found.

⁴ Use **INTO STRICT** to enforce a single-row query result. Otherwise, the “first” row is picked... 🐞

Assignment of Scalar Query Results

RHS of assignment $v := e$ is evaluated like a regular SQL query. In particular, e may be a *scalar subquery* in (...):

```
 $v := (Q)$   --  $Q$  yields single row, single column: 

|     |
|-----|
| $c$ |
|-----|


```

- Evaluates `SELECT (Q)` behind the scenes. Thus:
 - assigns cell value c cast to type τ of v , or
 - assigns `NULL :: τ` to v if Q returns no row, or
 - yields runtime error if Q returns more than one row or column (or if the cast fails). ⚠

6 | If All You Want Are the Side Effects...

1. Statement `NULL` does nothing (no side effects).
2. SQL **DML statements** (`INSERT/DELETE/UPDATE`) without `RETURNING` clauses are valid PL/SQL statements: no value is returned, the effect on the database is performed.
3. A SQL **query** `SELECT ...<query>...` may be performed solely for its side effects (e.g., invocation of a side-effecting UDF) as well:

PERFORM ...<query>... --  **PERFORM** replaces the **SELECT** keyword

Resulting rows are discarded (but variable `FOUND` is set).

7 : Returning From a Non-Table Function (RETURNS τ)

RETURN e

1. Evaluate e , cast value to return type τ of the function.
 - If $\tau \equiv \text{void}$, omit e . A `void` function whose control flow reaches the end of the top-level block, returns automatically.
2. Execution resumes in the calling function or query which receives the returned value.

To return multiple values, declare the function to return a row type.

“Returning” From a Table Function (**RETURNS SETOF τ**)

1 RETURN NEXT e ;
 s

2 RETURN QUERY Q ;
 s

- Add (bag semantics: \cup) to the result table computed by the function. Execution resumes with following statement s — no return to the caller yet.
 - ① Evaluate expression e , add scalar/row to result.
 - ② Evaluate SQL query Q , append all rows to result.
- Use plain **RETURN;** to return the result table accumulated so far and resume execution in the caller.

9 | Iterated Statements

```

1                                LOOP sG END LOOP
2                                WHILE p LOOP sG END LOOP
3    FOR vi IN [ REVERSE ] e0..e1 [ BY e2 ] LOOP sG END LOOP
4                                FOR vr IN q LOOP sG END LOOP
5    FOREACH va IN [ SLICE n ] ARRAY ea LOOP sG END LOOP

```

- ① Endless loop (see [EXIT](#) below).
- ② [p](#) :: [bool](#).
- ③ [e_{0,1,2}](#) :: [int](#). No [BY](#): [e₂](#) ≡ 1. [v_i](#) :: [int](#) (auto-declared) bound to [e₀](#), [e₀ ± 1×e₂](#), [e₀ ± 2×e₂](#), ... ([REVERSE](#): [±](#) ≡ [-](#)).
- ④ [q](#) SQL query. [v_r](#) successively bound to resulting rows.
- ⑤ [e_a](#) :: [τ](#)[]. No [SLICE](#): [v_a](#) :: [τ](#) bound to array elements.
[SLICE n](#): [v_a](#) :: [τ](#)[] bound to sub-arrays in [n](#)th dimension.

Leaving/Short-Cutting Loops

All five **LOOP** forms support optional `<< <label> >>` prefixes:

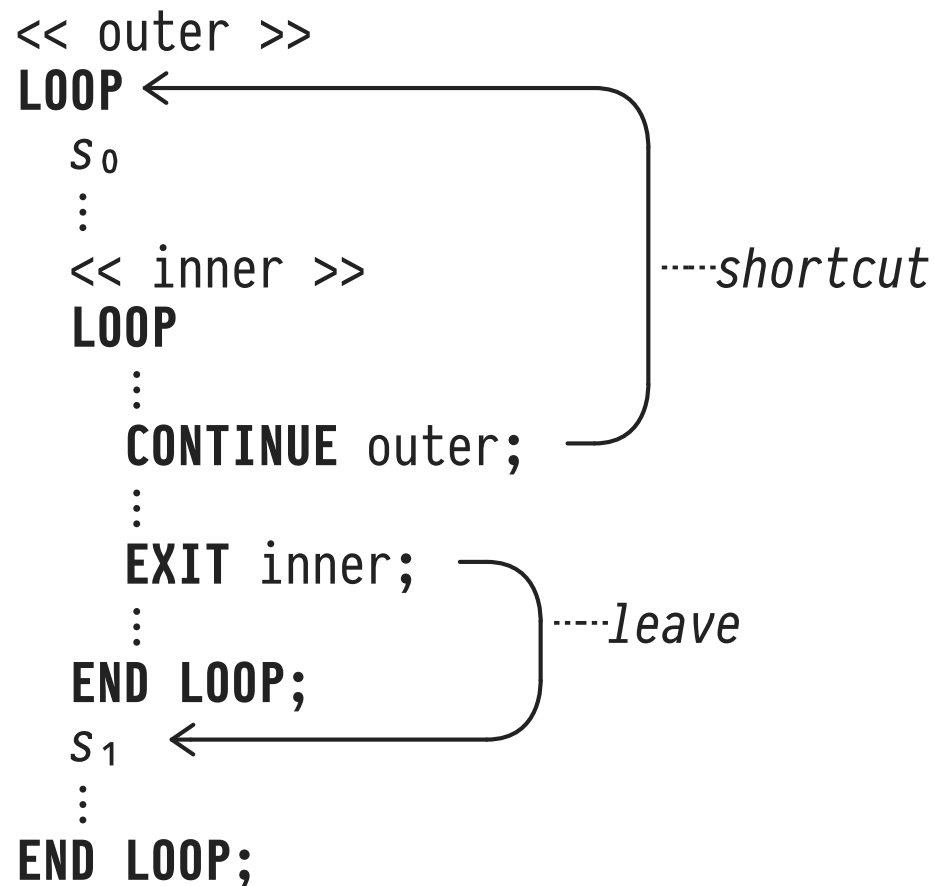
```
<< <label> >> ... LOOP sG END LOOP
```

We may alter the control flow inside a loop via:

```
1      EXIT [<label>] [ WHEN p ]  
2 CONTINUE [<label>] [ WHEN p ]
```

- No `<label>`: refer to innermost enclosing loop.
- **WHEN** *p*: leave/shortcut loop only if *p* \equiv `true`.
- **EXIT** `<label>` may also be used to leave a statement block.

Leaving/Shortcutting Loops



- Shortcutting a `WHILE p` loop leads to re-evaluation of *p*.

10 : Trapping Exceptions in Blocks

```

BEGIN
  :    -- } errors or RAISE ex statements transfer control
  sx -- } to the EXCEPTION clause – if sx changed the
  :    -- } database, also performs a rollback
EXCEPTION
  [ WHEN exi1 [, exij]* THEN si ]+
END;
s1    -- next statement if no exception occurred

```

- On error or **RAISE**, search for first matching exception category/name *ex_{ij}*, execute *s_i*, then *s₁*.
- If no match is found (or *s_i* fails), propagate exception to enclosing block. Abort function if in outermost block.

Raising Exceptions

one expression per '%' in message

```

❶ RAISE [ <level> ] '... % ... % ...' [, e]*
❷ RAISE [ <level> ] ex
❸ ASSERT p [, e]
  
```

- *level* ∈ {DEBUG, LOG, INFO, NOTICE, WARNING}. Only the default *level* = EXCEPTION raises an exception of name RAISE_EXCEPTION (or *ex*⁵, if provided).
- ASSERT *p* (*p* :: bool) raises exception ASSERT_FAILURE— with optional message *e* :: text—if *p* ≡ false.

⁵ See <https://www.postgresql.org/docs/current/errcodes-appendix.html> for a catalog of exception categories/names.